TAKE-HOME ASSIGNMENT

AI Generalist • Hiring Challenge

TIME BUDGET	DELIVERABLE	SUBMIT WITHIN	STACK
5–6 hours	**Working system + Loom**	**5 days of receipt**	**Anything goes**

CONTEXT

Confido Health is a voice AI company that automates administrative tasks for healthcare practices. We're a 50-person team growing fast. The People team is overwhelmed: every open role burns ~18 hours of manual work across screening, scheduling, interviewing, and onboarding. Most of it is repetitive judgment work that's perfect for AI — but nobody has built the system yet.

That's the job. As our AI Generalist, you'd parachute into teams like this, diagnose what's broken, and ship something that works. This assignment simulates exactly that.

THE CHALLENGE

Build an end-to-end hiring automation system that takes a job description and a folder of 20 resumes as input, and produces:

1. **A ranked shortlist** with per-candidate reasoning that a hiring manager can actually trust and act on (not just a similarity score).
2. **Auto-generated outreach** for shortlisted candidates — personalized emails referencing specific things from their resume, not generic templates.
3. **A structured interview kit** per shortlisted candidate — tailored questions that probe gaps or claims in their specific background, mapped to the role's evaluation criteria.

The hard part: We'll test your system with messy, real-world inputs — resumes in different formats (PDF, DOCX, even images of printed CVs), inconsistent naming, career gaps, job titles that don't match standard taxonomies, and a JD that's intentionally vague in places. Your system needs to handle all of this gracefully, not crash or silently produce garbage.

WHY THIS IS HARD

This isn't a wrapper around a single LLM call. To do this well, you'll need to think about:

- How to parse wildly inconsistent resume formats into a structured representation — and what to do when parsing fails.
- How to evaluate candidates against criteria that aren't fully specified in the JD (because they never are in the real world).
- How to produce reasoning that's specific enough to be useful, not just "strong match" / "weak match."
- How to generate emails that feel like a real recruiter wrote them, not an AI that read a LinkedIn profile.
- How to create interview questions that actually probe the interesting parts of each candidate's background.
- How to architect this so it's not a fragile script but something a non-technical person could eventually use.

EVALUATION CRITERIA

Criteria	Weight	What separates great from good
System Design	**30%**	You didn't just chain prompts — you designed a pipeline with clear stages, error handling, and fallbacks. Bonus: you thought about what happens when the LLM is wrong.
Output Quality	**25%**	The shortlist reasoning is specific and defensible. The emails don't read like AI slop. The interview questions reveal something a generic question wouldn't.
Robustness	**20%**	We'll feed in garbage and edge cases. Your system degrades gracefully — flags what it can't parse, doesn't hallucinate qualifications, handles the ugly stuff.
Judgment & Taste	**15%**	You made smart trade-offs about what to automate vs. leave to a human. You chose the right tools. Your prompts show craft, not just copy-paste.
Communication	**10%**	Your Loom tells a clear story: what you built, why, what's brittle, and what you'd do with more time. Your README is something we'd actually read.

SUBMISSION

- GitHub repo (public or invite us) with all code, prompts, and a README.
- Include sample input data (resumes + JD) and the outputs your system generated, so we can verify results.
- README must cover: architecture decisions, how to run it, known limitations, and what you'd build next.
- A 5–10 minute Loom: demo the system live, walk us through the code, and be honest about what's brittle.

WHAT WE DON'T CARE ABOUT

- Pretty UI — a CLI or notebook is fine. We're evaluating your AI engineering and product thinking, not your CSS.
- Which LLM provider you use. GPT-4, Claude, Gemini, local models, a mix — whatever works.
- Whether you used AI tools to build this. We expect you to. Just be honest about what you used and how.

Questions? *Reply to the email that sent you this. Asking smart clarifying questions is a positive signal, not a penalty.*